

# LSTM\_Model

May 7, 2021

```
[1]: import numpy as np
import pandas as pd
from ast import literal_eval
import time

from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedShuffleSplit

import tensorflow as tf
from tensorflow.keras.preprocessing.sequence import pad_sequences

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset, TensorDataset
from torch.nn.utils.rnn import pack_padded_sequence, pad_packed_sequence

DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
torch.backends.cudnn.deterministic = True
```

## 0.1 Import and Pad Data

```
[2]: ## Import Data

# sequence column lists are in string format so use literal eval to reconvert
blk_msg_sequence = pd.read_csv('Data/blk_event_sequences.csv', header=0, index_col=None, converters={"sequence": literal_eval})

# Pad sequences (Predictor)
padded_seq = pad_sequences(blk_msg_sequence.sequence, padding='post')

# Anomaly Labels (Response)
labels = blk_msg_sequence['anomaly']

#print(f'Sequence Example: {padded_seq[19]}')
```

```

#print(f'Label Example: {labels[19]}')
print(f'Percent Anomalous: {round(sum(labels)/len(labels)*100, 2)} %')

# all_events = blk_msg_sequence.sequence.apply(pd.Series).stack().
  ↳reset_index(drop = True)
# Unique Words (75)
# len(all_events.unique())
# > 75

```

Percent Anomalous: 2.93 %

## 0.2 Split Data

```

[3]: # Split Train 0.70
sss = StratifiedShuffleSplit(n_splits=1, train_size=0.70)
for train_index, rest_index in sss.split(padded_seq, labels):
    x_train, x_other = padded_seq[train_index], padded_seq[rest_index]
    y_train, y_other = labels[train_index], labels[rest_index]

# Split Val 0.10 / Test 0.20
x_val, x_test, y_val, y_test = train_test_split(x_other, y_other, test_size=2/3)

# Standardize to arrays
y_train, y_val, y_test = np.array(y_train), np.array(y_val), np.array(y_test)

# Check proportions
trn, val, tst = len(y_train), len(y_val), len(y_test)
tot = sum([trn, val, tst])
print(f'Train: {round(trn/tot, 2)} ({trn})',
      f'Val: {round(val/tot, 2)} ({val})',
      f'Test: {round(tst/tot, 2)} ({tst})',
      f'Total: {1} ({tot})', sep = '\n')

```

Train: 0.7 (402542)  
 Val: 0.1 (57506)  
 Test: 0.2(115013)  
 Total: 1 (575061)

## 0.3 Build Data Loaders

```

[4]: BATCH_SIZE = 1000

#####
## Create Tensor Datasets

```

```
#####
train_data = TensorDataset(torch.from_numpy(x_train), torch.from_numpy(y_train))
val_data = TensorDataset(torch.from_numpy(x_val), torch.from_numpy(y_val))
test_data = TensorDataset(torch.from_numpy(x_test), torch.from_numpy(y_test))

#####
## Dataset Iterators:
#####
train_loader = DataLoader(train_data, shuffle=True, batch_size=BATCH_SIZE,
    ↪pin_memory=True, drop_last=True)
val_loader = DataLoader(val_data, shuffle=True, batch_size=BATCH_SIZE,
    ↪pin_memory=True, drop_last=True)
test_loader = DataLoader(test_data, shuffle=True, batch_size=BATCH_SIZE,
    ↪pin_memory=True, drop_last=True)

print('Training Dimensions\n')
for data_cat in [train_loader, val_loader, test_loader]:
    for batch in data_cat:
        print(f'Sequence matrix size: {batch[0].size()}')
        print(f'Target vector size: {batch[1].size()}\n')
        break
```

Training Dimensions

Sequence matrix size: torch.Size([1000, 298])  
 Target vector size: torch.Size([1000])

Sequence matrix size: torch.Size([1000, 298])  
 Target vector size: torch.Size([1000])

Sequence matrix size: torch.Size([1000, 298])  
 Target vector size: torch.Size([1000])

## 0.4 PyTorch LSTM Model

And associated functions \*\*\*

```
[5]: #####
## ARCHITECTURAL PARAMETERS
#####
VOCAB_SIZE = 75+1
OUTPUT_SIZE = 1
EMBEDDING_DIM = 64
HIDDEN_DIM = 128
CLIP = 5
```

```
#####
## MODEL
#####
class Anomaly_Detector(nn.Module):
    def __init__(self, vocab_size, output_size, embedding_dim, hidden_dim,
        ↪ batch_size, n_layers, bidirectinoal=True, drop_prob=0):
        super().__init__()

        # Params
        self.batch_size = batch_size
        self.n_layers = n_layers
        self.bidir = 2 if bidirectinoal==True else 1
        self.hidden_dim = hidden_dim

        # Layers
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.lstm = nn.LSTM(embedding_dim, hidden_dim,
            batch_first=True,
            num_layers=n_layers,
            dropout=drop_prob,
            bidirectional=True)

        self.fc = nn.Linear(hidden_dim*self.bidir, output_size) # multiply
        ↪ hidden by 2 if using bidirectinoality
        self.activation = nn.Sigmoid()
        self.dropout = nn.Dropout(drop_prob)

    def forward(self, sequences):
        sequences = sequences.long().to(DEVICE) # [batch_size, seq_len]
        embedded = self.dropout(self.embedding(sequences)) # [batch_size,
        ↪ seq_length, embedded_dim]
        lstm_out, (hidden, cell) = self.lstm(embedded)
        # lstm_out = [batch_size, seq_length, hidden_dim * num_directions]
        # hidden = [num layers * num directions, batch size, hidden_dim]
        # cell = [num layers * num directions, batch size, hidden_dim]

        if self.bidir==2:
            # Concat the final forward and barckward hidden layers + apply
            ↪ dropout if specified
            hidden = self.dropout(torch.cat((hidden[-2,:,:], hidden[-1,:,:]),
            ↪ dim = 1))
        elif self.bidir==1:
            hidden = self.dropout(hidden[-1,:,:])

        out = self.activation(self.fc(hidden).reshape(-1))

        return out
```

```

#####
## DEFINE ACCURACY FUNCTION
#####
def compute_binary_accuracy(model, data_loader, device, print_output=False):
    model.eval()
    all_preds = torch.empty(0).cuda()
    all_labs = torch.empty(0).cuda()

    with torch.no_grad():
        for inpt, labs in data_loader:

            inpt, labs = inpt.cuda(), labs.cuda()
            all_labs = torch.cat((all_labs, labs), 0).reshape(-1)

            logits = model(inpt)
            predicted_labels = (logits > 0.5).reshape(-1)
            all_preds = torch.cat((all_preds, predicted_labels), 0).reshape(-1)

        # Move to CPU
        all_preds, all_labs = all_preds.cpu(), all_labs.cpu()
        accuracy = round(accuracy_score(all_labs, all_preds), 4) # of all
        ↳ predictions, how many were correct?
        precision = round(precision_score(all_labs, all_preds), 4) # of
        ↳ predicted anomalies, how many were correct?
        recall = round(recall_score(all_labs, all_preds), 4) # of all true
        ↳ anomalies, how many were recovered?

        if print_output:
            print(confusion_matrix(all_labs, all_preds), end = '\n\n')
            print('Accuracy:', accuracy)
            print('Precision:', precision)
            print('Recall:', recall)
            return None

    return accuracy, precision, recall

#####
## DEFINE TRAINING + TESTING FUNCTION
#####
def training_function():
    print('\nSTARTING TRAINING\n')
    start_time = time.time()

    for epoch in range(NUM_EPOCHS):

```

```

model.train()
print(f'Current Epoch: {epoch+1}')

for batch_idx, batch_data in enumerate(train_loader):

    inpt, labs = batch_data[0], batch_data[1].to(DEVICE)

    optimizer.zero_grad() # Clear gradients from last batch

    logits = model(inpt) # Compute logits (ie. final hidden state)

    # Compute cost
    cost = criterion(logits.squeeze(), labs.float())
    cost.backward()

    # Use clip_grad_norm to prevent exploding gradients
    nn.utils.clip_grad_norm_(model.parameters(), CLIP)
    optimizer.step()

    # Log
    if not batch_idx % 100:
        print (f'Batch {batch_idx:03d}/{len(train_loader):03d} | '
              f'BCE Cost: {cost:.4f}')

    with torch.set_grad_enabled(False):
        model.eval()
        print(f'Epoch Comleted: {epoch+1}'
              f'\n\twith training accuracy: {compute_binary_accuracy(model,
→train_loader, DEVICE)[0]:.4f}%')
        print(f'\n\tand validation accuracy:
→{compute_binary_accuracy(model, val_loader, DEVICE)[0]:.4f}%')

        print(f'\tTime elapsed: {(time.time() - start_time)/60:.2f} min\n')

    ## Final Test Accuracy
    print(f'Total Training Time: {(time.time() - start_time)/60:.2f} min')
    print('-----')
    print('Final Test Accuracy/Precision/Recall : ',
→compute_binary_accuracy(model, test_loader, DEVICE))
    print('-----')

```

```

[7]: # Make models reproducible
RANDOM_SEED = 123

```

```

[7]: # Params: VANILLA

torch.manual_seed(RANDOM_SEED)

```

```

LEARNING_RATE = 0.0001
NUM_EPOCHS = 3
NUM_LAYERS = 1
DROPOUT = 0
BIDIRECTIONAL = False

## INITIALIZE MODEL
model = Anomaly_Detector(VOCAB_SIZE, OUTPUT_SIZE, EMBEDDING_DIM, HIDDEN_DIM,
    ↪BATCH_SIZE, NUM_LAYERS, BIDIRECTIONAL, DROPOUT)
model = model.to(DEVICE)

## LOSS AND OPTIMIZER
criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

## TRAIN/TEST
training_function()

```

STARTING TRAINING

Current Epoch: 1

Batch 000/402 | BCE Cost: 0.6704

Batch 100/402 | BCE Cost: 0.1283

Batch 200/402 | BCE Cost: 0.0858

Batch 300/402 | BCE Cost: 0.0890

Batch 400/402 | BCE Cost: 0.1065

Epoch Completed: 1

with training accuracy: 0.9815%

and validation accuracy: 0.9816%

Time elapsed: 0.85 min

Current Epoch: 2

Batch 000/402 | BCE Cost: 0.0928

Batch 100/402 | BCE Cost: 0.0595

Batch 200/402 | BCE Cost: 0.0224

Batch 300/402 | BCE Cost: 0.0107

Batch 400/402 | BCE Cost: 0.0055

Epoch Completed: 2

with training accuracy: 0.9981%

and validation accuracy: 0.9980%

Time elapsed: 1.65 min

Current Epoch: 3

Batch 000/402 | BCE Cost: 0.0112

Batch 100/402 | BCE Cost: 0.0074

Batch 200/402 | BCE Cost: 0.0180

Batch 300/402 | BCE Cost: 0.0044

Batch 400/402 | BCE Cost: 0.0152  
Epoch Completed: 3  
    with training accuracy: 0.9989%  
    and validation accuracy: 0.9987%  
    Time elapsed: 2.45 min

Total Training Time: 2.45 min

-----  
Final Test Accuracy/Precision/Recall : (0.9989, 0.9798, 0.9827)  
-----

```
[7]: # Params: 2 Layer

torch.manual_seed(RANDOM_SEED)
LEARNING_RATE = 0.0001
NUM_EPOCHS = 3
NUM_LAYERS = 2
DROPOUT = 0
BIDIRECTIONAL = False

## INITIALIZE MODEL
model = Anomaly_Detector(VOCAB_SIZE, OUTPUT_SIZE, EMBEDDING_DIM, HIDDEN_DIM,
    ↪BATCH_SIZE, NUM_LAYERS, BIDIRECTIONAL, DROPOUT)
model = model.to(DEVICE)

## LOSS AND OPTIMIZER
criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

training_function()
```

STARTING TRAINING

Current Epoch: 1  
Batch 000/402 | BCE Cost: 0.6993  
Batch 100/402 | BCE Cost: 0.0863  
Batch 200/402 | BCE Cost: 0.1178  
Batch 300/402 | BCE Cost: 0.0860  
Batch 400/402 | BCE Cost: 0.0934  
Epoch Completed: 1  
    with training accuracy: 0.9815%  
    and validation accuracy: 0.9808%  
    Time elapsed: 1.85 min

Current Epoch: 2  
Batch 000/402 | BCE Cost: 0.0537  
Batch 100/402 | BCE Cost: 0.0155



Batch 200/402 | BCE Cost: 0.0071  
Batch 300/402 | BCE Cost: 0.0072  
Batch 400/402 | BCE Cost: 0.0089  
Epoch Completed: 2  
    with training accuracy: 0.9981%  
    and validation accuracy: 0.9982%  
    Time elapsed: 3.88 min

Current Epoch: 3  
Batch 000/402 | BCE Cost: 0.0061  
Batch 100/402 | BCE Cost: 0.0168  
Batch 200/402 | BCE Cost: 0.0101  
Batch 300/402 | BCE Cost: 0.0088  
Batch 400/402 | BCE Cost: 0.0032  
Epoch Completed: 3  
    with training accuracy: 0.9989%  
    and validation accuracy: 0.9990%  
    Time elapsed: 5.93 min

Total Training Time: 5.93 min

-----  
Final Test Accuracy/Precision/Recall : (0.9989, 0.9801, 0.9813)  
-----

```
[8]: # Params 1 Layer / Bidirectional

torch.manual_seed(RANDOM_SEED)
LEARNING_RATE = 0.0001
NUM_EPOCHS = 3
NUM_LAYERS = 1
DROPOUT = 0
BIDIRECTIONAL = True

## INITIALIZE MODEL
model = Anomaly_Detector(VOCAB_SIZE, OUTPUT_SIZE, EMBEDDING_DIM, HIDDEN_DIM,
    ↪BATCH_SIZE, NUM_LAYERS, BIDIRECTIONAL, DROPOUT)
model = model.to(DEVICE)

## LOSS AND OPTIMIZER
criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

training_function()
```

STARTING TRAINING

Current Epoch: 1

```
Batch 000/402 | BCE Cost: 0.7502
Batch 100/402 | BCE Cost: 0.1390
Batch 200/402 | BCE Cost: 0.1049
Batch 300/402 | BCE Cost: 0.0787
Batch 400/402 | BCE Cost: 0.0833
Epoch Comleted: 1
    with training accuracy: 0.9815%
    and validation accuracy: 0.9809%
    Time elapsed: 0.90 min
```

```
Current Epoch: 2
Batch 000/402 | BCE Cost: 0.1048
Batch 100/402 | BCE Cost: 0.0760
Batch 200/402 | BCE Cost: 0.0169
Batch 300/402 | BCE Cost: 0.0080
Batch 400/402 | BCE Cost: 0.0135
Epoch Comleted: 2
    with training accuracy: 0.9980%
    and validation accuracy: 0.9982%
    Time elapsed: 1.79 min
```

```
Current Epoch: 3
Batch 000/402 | BCE Cost: 0.0170
Batch 100/402 | BCE Cost: 0.0036
Batch 200/402 | BCE Cost: 0.0027
Batch 300/402 | BCE Cost: 0.0169
Batch 400/402 | BCE Cost: 0.0020
Epoch Comleted: 3
    with training accuracy: 0.9991%
    and validation accuracy: 0.9991%
    Time elapsed: 2.68 min
```

```
Total Training Time: 2.68 min
```

```
-----
Final Test Accuracy/Precision/Recall : (0.9991, 0.993, 0.9777)
-----
```

```
[8]: # Params: 2 layer / Bidirectional
```

```
torch.manual_seed(RANDOM_SEED)
LEARNING_RATE = 0.0001
NUM_EPOCHS = 3
NUM_LAYERS = 2
DROPOUT = 0
BIDIRECTIONAL = True

## INITIALIZE MODEL
```

```

model = Anomaly_Detector(VOCAB_SIZE, OUTPUT_SIZE, EMBEDDING_DIM, HIDDEN_DIM,
↪BATCH_SIZE, NUM_LAYERS, BIDIRECTIONAL, DROPOUT)
model = model.to(DEVICE)

## LOSS AND OPTIMIZER
criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

training_function()

## See confusion matrix
compute_binary_accuracy(model, test_loader, DEVICE, print_output=True)

```

## STARTING TRAINING

Current Epoch: 1

Batch 000/402 | BCE Cost: 0.7074

Batch 100/402 | BCE Cost: 0.0932

Batch 200/402 | BCE Cost: 0.1304

Batch 300/402 | BCE Cost: 0.0753

Batch 400/402 | BCE Cost: 0.0239

Epoch Completed: 1

with training accuracy: 0.9816%

and validation accuracy: 0.9813%

Time elapsed: 2.35 min

Current Epoch: 2

Batch 000/402 | BCE Cost: 0.0371

Batch 100/402 | BCE Cost: 0.0210

Batch 200/402 | BCE Cost: 0.0154

Batch 300/402 | BCE Cost: 0.0086

Batch 400/402 | BCE Cost: 0.0021

Epoch Completed: 2

with training accuracy: 0.9983%

and validation accuracy: 0.9982%

Time elapsed: 4.21 min

Current Epoch: 3

Batch 000/402 | BCE Cost: 0.0077

Batch 100/402 | BCE Cost: 0.0051

Batch 200/402 | BCE Cost: 0.0034

Batch 300/402 | BCE Cost: 0.0010

Batch 400/402 | BCE Cost: 0.0008

Epoch Completed: 3

with training accuracy: 0.9995%

and validation accuracy: 0.9996%

Time elapsed: 6.05 min

Total Training Time: 6.05 min

-----  
Final Test Accuracy/Precision/Recall : (0.9995, 0.9943, 0.9895)  
-----

[[111632 19]  
[ 35 3314]]

Accuracy: 0.9995

Precision: 0.9943

Recall: 0.9895

[ ]: